

1 Intro

This document is concerned with the formal semantics of probabilistic programs used for probabilistic modeling and machine learning. One approach in this setting is that

- we want programs in our language to specify probabilistic generative models of data, and
- we want these programs to be interpreted as (un-normalized) posterior distributions over variables of interest in the model (latents and model parameters)

That is, we want our compilers/interpreters to do probabilistic inference. Once one specifies a programming language, including its terms, types, and syntax, one must specify how to interpret what a program means: in our case we must specify how a generative model maps to a posterior distribution. We will focus on denotational semantics, which associate a mathematical object with each fragment of program code in a compositional way. As such, each piece of our program will describe a distribution and the whole program will too.

We present a walkthrough of *Semantics for probabilistic programming: higher-order functions, continuous distributions, and soft constraints* [1, 2016] and *Commutative semantics for probabilistic programming* [2, 2017], mainly focusing on the presentation in [2]. We first briefly review measures and kernels because they show up in the semantics. We will later be interested in approximate inference but focus on the exact case here.

2 Measure Theory

A **σ -algebra** Σ_X on a set X is a collection of subsets of X that contains \emptyset and is closed under complements and countable unions. A **measurable space** is a pair (X, Σ_X) of a set and a σ -algebra on it. The elements of Σ_X , which are themselves sets of elements in X , are called **measurable sets**. For example, the Borel sets are the smallest σ -algebra on \mathbb{R} that contains the intervals. This is the usual σ -algebra for \mathbb{R} . For any countable set, you can always start off with the individual elements of the set: satisfying complements and unions this gives you the powerset σ -algebra.

A **measure** on a measurable space (X, Σ_X) is a function $\mu : \Sigma_X \rightarrow [0, \infty]$ into the set $[0, \infty]$ of extended non-negative reals that takes countable disjoint unions to sums, i.e. $\mu(\emptyset) = 0$ and $\mu(\bigcup_{n \in \mathbb{N}} U_n) = \sum_{n \in \mathbb{N}} \mu(U_n)$ for any \mathbb{N} -indexed sequence of disjoint measurable sets U_n . A **probability measure** is a measure μ such that $\mu(X) = 1$. For example, the Lebesgue measure λ on \mathbb{R} is generated by $\lambda((a, b)) = b - a$. For any $x \in X$, the Dirac measure $\delta_x(U) = 1$ if $x \in U$ and 0 otherwise.

A **kernel** k from X to Y , notated $k : X \rightsquigarrow Y$, is a function $k : X \times \Sigma_Y \rightarrow [0, \infty]$ such that

- for fixed x , $k(x, -) : \Sigma_Y \rightarrow [0, \infty]$ is a measure
- for fixed dy , $k(-, dy) : X \rightarrow [0, \infty]$ is a measurable function.

An intuitive example of kernels at work is conditional probability. Take $k(x, -)$ to be a probability measure $\mu_Y : \Sigma_Y \rightarrow [0, 1]$ on Y given a particular value $X = x$: $\sum_{dy \in \Sigma_y} k(x, dy) = 1$. Note that $\sum_{x \in X} k(x, dy) \neq 1$ in general.

3 Types and Semantics

The two papers [1,2] don't actually define the set of terms of the language, but they are implied to be the typical ones in a language like Haskell. They just define a set of types and typing rules. The 2018 POPL paper on semantics for higher-order languages [3] does formally introduce a kind and type system and a set of terms. We define our types as:

$$\mathbb{A}, \mathbb{B} ::= \mathbb{R} | P(\mathbb{A}) | 1 | \mathbb{A} \times \mathbb{B} | \sum_i \mathbb{A}_{i \in I}$$

where I is countable and non-empty. As we will see, types are to be interpreted as measurable spaces $[[\mathbb{A}]]$. We have sum and product types. We have a type \mathbb{R} that denotes the Reals $[[\mathbb{R}]]$. We have a type $P(\mathbb{A})$ that denotes the set of probability measures $[[P(\mathbb{A})]]$ over a space denoted by \mathbb{A} . We have the type 1 that denotes a singleton set. For booleans, we can just take $(1 + 1)$ and $P(1 + 1)$ is the type for distributions over booleans. For the natural numbers, we can use $\sum_{i \in \mathbb{N}} 1$. Our terms are therefore elements of these sets, like elements of a σ -algebra whose probability we are interested in and density functions.

Typing judgements: $\Gamma \vdash_d$ for deterministic judgements and $\Gamma \vdash_p$ for probabilistic judgements. Having the two typing judgements is mainly for notational clarity: it helps us to define interpretation differently for deterministic and probabilistic terms.

3.1 Typing Judgements and Interpretations for Deterministic Terms

Deterministic terms $\Gamma \vdash_d t : \mathbb{A}$ are interpreted as measurable functions $[[t]] : [[\Gamma]] \rightarrow [[\mathbb{A}]]$, where $[[t]](\gamma) = x$ is an element of the underlying set $[[\mathbb{A}]]$ and not a measurable set in $\Sigma_{[[\mathbb{A}]]}$. **Note:** the following informally uses the notation $\Gamma :: (x : \mathbb{A}) :: \Gamma'$ to show a particular variable x of type \mathbb{A} in the environment.

- **basic term** $x : \mathbb{A}$

$$\frac{}{\Gamma :: (x : \mathbb{A}) :: \Gamma' \vdash_d x : \mathbb{A}} \quad [[x]](\gamma :: d :: \gamma') = d$$

- **Disjoint Sum Type**

$$\frac{\Gamma \vdash_d t : \mathbb{A}_i}{\Gamma \vdash_d (i, t) : \sum_{i \in I} \mathbb{A}_i} \quad [[(i, t)]](\gamma) = (i, [[t]](\gamma))$$

- **Deterministic case**

$$\frac{\Gamma \vdash_d t : \sum_{i \in I} \mathbb{A}_i \quad (\Gamma :: (x : \mathbb{A}_i) \vdash_d u_i : \mathbb{B})_{i \in I}}{\Gamma \vdash_d (\text{case } t \text{ of } \{(i, x) \rightarrow u_i\}_{i \in I}) : \mathbb{B}}$$

$$[[\text{case } t \text{ of } \{(i, x) \rightarrow u_i\}_{i \in I}]](\gamma) = [[u_i]](\gamma, d) \text{ if } [[t]](\gamma) = (i, d)$$

- **Unit ()**

$$\frac{}{\Gamma \vdash_d () : 1} \quad [[()]](\gamma) = ()$$

- **Product Type**

$$\frac{\Gamma \vdash_d t_0 : \mathbb{A}_0 \quad \Gamma \vdash_d t_1 : \mathbb{A}_1}{\Gamma \vdash_d (t_0, t_1) : \mathbb{A}_0 \times \mathbb{A}_1} \quad [[(t_0, t_1)]](\gamma) = ([[t_0]](\gamma), [[t_1]](\gamma))$$

- **Projection**

$$\frac{\Gamma \vdash_d t : \mathbb{A}_0 \times \mathbb{A}_1}{\Gamma \vdash_d \pi_j(t) : \mathbb{A}_j} \quad [[\pi_j(t)]](\gamma) = d_j \text{ if } [[t]](\gamma) = (d_0, d_1)$$

where the **case** expression above has a deterministic continuation u_i . We will come back to the probabilistic continuation case. Now the semantics for sequencing.

3.2 Typing Judgements and Interpretations for Probabilistic Terms

Probabilistic terms $\Gamma \vdash_p t : \mathbb{A}$ are interpreted as s-finite kernels $[[t]] : [[\Gamma]] \rightsquigarrow [[\mathbb{A}]]$.

- **return(t)**

$$\frac{\Gamma \vdash_d t : \mathbb{A}}{\Gamma \vdash_p \text{return}(t) : \mathbb{A}} \quad [[\text{return}(t)]](\gamma, da) = \delta_{[[t]](\gamma)}(da)$$

this corresponds to a Dirac Delta measure sitting at a point $[[t]](\gamma)$.

- **let x = t in u**

$$\frac{\Gamma \vdash_p t : \mathbb{A} \quad \Gamma, x : \mathbb{A} \vdash_p u : \mathbb{B}}{\Gamma \vdash_p \text{let } x = t \text{ in } u : \mathbb{B}} \quad [[\text{let } x = t \text{ in } u]](\gamma, db) = \int_{x, dx \in [[\mathbb{A}]]} [[u]](\gamma :: x, db) [[t]](\gamma, dx)$$

- **probabilistic case**

$$\frac{\Gamma \vdash_d t : \sum_{i \in I} \mathbb{A}_i \quad (\Gamma, x : \mathbb{A}_i \vdash_p u_i : \mathbb{B})_{i \in I}}{\Gamma \vdash_p (\text{case } t \text{ of } \{(i, x) \rightarrow u_i\}_{i \in I}) : \mathbb{B}}$$

$$[[\text{case } t \text{ of } \{(i, x) \rightarrow u_i\}_{i \in I}]](\gamma, db) = [[u_i]](\gamma :: d, db) \text{ if } [[t]](\gamma) = (i, d)$$

- **sample(t)**

$$\frac{\Gamma \vdash_d t : P(\mathbb{A})}{\Gamma \vdash_p \text{sample}(t) : \mathbb{A}} \quad [[\text{sample}(t)]](\gamma, da) = \left([[t]](\gamma) \right)(da)$$

- **score(t)**

$$\frac{\Gamma \vdash_d t : \mathbb{R}}{\Gamma \vdash_p \text{score}(t) : 1}$$

$$[[\text{score}(t)]](\gamma, du) = \begin{cases} \text{abs}([t]](\gamma)), & \text{if } du = \{()\} \\ 0, & \text{if } du = \emptyset \end{cases}$$

3.3 Typing Judgements and Interpretation for Normalize TODO

4 Example Programs

Beta-Bernoulli model: in the following, we derive that “the un-normalized posterior is a measure defined by integrating the likelihood with respect to the prior”.

$$\begin{aligned}
& [[\text{let } x = \text{sample}(\text{Beta}(2, 2)) \text{ in } \text{score}(\text{Bernoulli}(1; x)); \text{return}(x)]](db) \\
&= \int_x [[\text{let } a = \text{score}(\text{Bernoulli}(1; x)) \text{ in } \text{return}(x)]](\gamma :: x, db) \quad [[\text{sample}(\text{Beta}(2, 2))]](\gamma, dx) \\
&= \int_x [[\text{let } a = \text{score}(\text{Bernoulli}(1; x)) \text{ in } \text{return}(x)]](\gamma :: x, db) \quad \text{Beta}(dx; 2, 2) \\
&= \int_x \left(\sum_{s \in \{\emptyset, \{()\}} [[\text{return}(x)]](\gamma :: x :: a, db) \quad [[\text{score}(\text{Bernoulli}(1; x))]](\gamma :: x, s) \right) \quad \text{Beta}(dx; 2, 2) \\
&= \int_x \left([[\text{return}(x)]](\gamma :: x :: a, db) \quad [[\text{score}(\text{Bernoulli}(1; x))]](\gamma :: x, \{()\}) \right) \quad \text{Beta}(dx; 2, 2) \\
&= \int_x \left([[\text{return}(x)]](\gamma :: x :: a, db) \quad \text{Bernoulli}(1; x) \right) \quad \text{Beta}(dx; 2, 2) \\
&= \int_x \left(\mathbb{1}[x \in db] \quad \text{Bernoulli}(1; x) \right) \quad \text{Beta}(dx; 2, 2) \\
&= \text{UnnormalizedPosterior}(db) = \int_{x \in db} \text{Bernoulli}(1; x) \quad \text{Beta}(dx; 2, 2) \\
&= \int_{x \in db} \left(x^{2+1-1} (1-x)^{2-1} \right)
\end{aligned}$$

This gives the expected result. We were able to interpret the program as an un-normalized posterior for a Beta-Bernoulli model.

Note on conjugacy: In this example, we consider a Beta-Bernoulli model. The Beta distribution is a continuous distribution over $[0, 1]$ whose shape is determined by two parameters α and β so that the density is $\text{Beta}(x; \alpha, \beta) = \frac{1}{C(\alpha, \beta)} x^{\alpha-1} (1-x)^{\beta-1}$ where $C(\alpha, \beta)$ is just a normalizing constant. The Bernoulli distribution is a discrete distribution over $\{0, 1\}$ with parameter $x \in [0, 1]$ defined by mass function $\text{Bernoulli}(d; x) = x$ if $d = 1$ and $1-x$ if $d = 0$. The two distributions have the following property: if x is distributed $\text{Beta}(\alpha, \beta)$ and you observe data-points $d_i \in \{0, 1\}$ with likelihood distribution $\prod_i \text{Bernoulli}(d_i; x)$ (Bernoulli with parameter x), then the posterior distribution over x is $\text{Beta}(\alpha + \#1, \beta + \#0)$, where there are $\#1$ 1's and $\#0$ 0's in the data. This is a result coming from the Beta and Bernoulli being a **conjugate pair**. If we were to wrap the entire program in a call to `normalize()`, we should get a measure with density $\text{Beta}(2 + 1, 2)$

TODO: Definition of Density: if $f(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}x^2}$, then: ...

TODO: For probability measure p with density g , recover the importance sampling algorithm for sampling from p by sampling from a Gaussian by showing:

$$\begin{aligned}
[[\text{sample}(p)]] &= [[\text{let } x = \text{lebesgue} \text{ in } \text{score}(p(x)); \text{return}(x)]] \\
&= [[\text{let } x = \text{gauss}(0, 1) \text{ in } \text{score}(1/f(x)); \text{score}(g(x)); \text{return}(x)]] \\
&= [[\text{let } x = \text{gauss}(0, 1) \text{ in } \text{score}(g(x)/f(x)); \text{return}(x)]]
\end{aligned}$$

5 Higher Order

...

6 Citation

1. Staton et al. Semantics for probabilistic programming: higher-order functions, continuous distributions, and soft constraints. LICS 2016. <https://arxiv.org/pdf/1601.04943.pdf>
2. Staton. Commutative semantics for probabilistic programming. ESOP 2017. ‘ <http://www.cs.ox.ac.uk/people/samuel.staton/papers/esop2017.pdf>
3. Scibior et al. Denotational validation of Bayesian inference. POPL 2018. <https://arxiv.org/pdf/1711.03219.pdf>